MDPI

Article

Governing Procured AI: What Public Records Establish about Supplier Oversight and Responses

Abstract

Public AI procurement disclosures and contracts identify suppliers and oversight provisions, but do not necessarily establish what actions occurred or their effects. This purposive critical documentary review examines three government audits and an implementation-team report (28 source-linked descriptions across nine programme families), alongside seven United Kingdom documents (14 attributed interpretations) and two retrospective cross-record checks. The sources describe recurring buyer monitoring, a tested supplier software rollout, an audit-linked data-retention correction and collaborative adjustments; none establishes that an exercised contractual remedy caused improved public-service outcomes. Two cross-record checks illustrate attribution hazards. A Newcastle AI-use disclosure and a related contact-centre award identify different registered companies (09162820 and FC034225). A North Northamptonshire inspection report published in March 2026 concerns an assessment ending before a September 2025 contract extension; publication timing cannot assign undated staff testimony to the later agreement. The published agreement also distinguishes restricted performance-monitoring access from permitted data-processing compliance audits, without demonstrating an exercised remedy. These cases demonstrate why supplier identity, observation date, oversight authority, response and outcome must remain distinct. Purposive selection and AI-assisted interpretation preclude prevalence estimates, causal attribution and superiority claims over existing assurance frameworks.

Keywords: public-sector AI; AI procurement; supplier accountability; post-deployment oversight; critical review; documentary analysis

Key Points for Practitioners

- Confirm the registered contracting entity, buyer, AI component and contract period before linking an AI-use disclosure to a procurement award.
- Record the period in which a service was observed or assessed separately from the date on which a report was published.
- Distinguish audit rights for data-processing compliance from permissions to monitor or benchmark AI-system performance.
- Preserve dated oversight requests, supplier responses and outcome measurements as linked but separately verifiable records.

Received:
Accepted:
Published:

Copyright: © 2026 by the authors. Submitted to *Soc. Sci.* for possible open access publication under the terms and conditions of the Creative Commons Attribution (CC BY) license.

1. Introduction

Public organisations that acquire artificial intelligence (AI) services remain responsible for public decisions, even when suppliers control software releases, evaluation interfaces, hosting and technical records. Procurement therefore concerns both acquiring a capability and establishing the conditions under which its operation can be scrutinised. Contractual

approaches to algorithmic accountability have long addressed transparency, audits, privacy and participation; lifecycle oversight is also explicit in public procurement guidance (Coglianese and Lampmann 2021; Government 2020). The question is not whether these objectives are desirable, but which acquisition, oversight and response claims are supported by the public records examined.

Three claims require different evidence: an agency has a right to investigate a supplier; oversight produced an identifiable response; and that response improved a specified outcome. A signed contractual provision can support the first claim without establishing the second. An observed correction can support the second without estimating the third. Conversely, a correction induced by statutory inspection may be important even when no exercised contractual clause is visible. Treating contract-enforced causal benefit as the sole acceptable endpoint would obscure these other mechanisms.

This review asks: **What do selected public records establish about AI supplier oversight and responses during use, and where do acquisition, source-attribution and outcome boundaries limit stronger conclusions?** Two subsidiary questions concern the acquisition boundary—whether the tool was actually procured—and the relationship between document authority and the evidential status of a statement. The review focuses on AI as an acquired capability, rather than AI used to automate procurement administration. Here, oversight means a reported or documented review, test, information request or scrutiny episode involving a public buyer and a supplier-linked capability. A supplier response is an identifiable action attributed to the supplier; it is not synonymous with a contractual remedy or a beneficial effect. Routine monitoring and collaborative improvement are described as accountability-relevant activities. Neither establishes, by itself, that the supplier was held accountable or that service users benefited.

The contribution is a bounded comparative synthesis rather than a new assurance framework. It has three parts. First, it distinguishes acquisition records from AI-component contracts and separates documented reviews or corrections from the exercise of contractual powers and independently measured effects. Second, it illustrates two source-specific cross-record linkage hazards: mismatched legal-entity identifiers and publication dates that postdate the observations they report. Third, it examines a published agreement in which a data-processing-compliance audit provision and a separate restriction on performance-monitoring access cover different objects of scrutiny; neither provision alone establishes an unqualified AI-performance evaluation right. These comparisons acknowledge documented governance activity without implying contractual compulsion, technical improvement or public benefit. They do not establish that procurement contracts generally succeed or fail.

2. Existing research and the contribution boundary

Research on AI procurement already makes a strong normative and institutional case for controlling supplier relationships. Coglianese and Lampmann explain how contractual terms can address accountability without waiting for new legislation (Coglianese and Lampmann 2021). Sánchez-Graells examines the legal and institutional difficulties of acquiring AI within public administration (S'anchez-Graells 2024). These contributions establish why the contracting process matters; they do not make every procurement record evidence of an exercised remedy.

Empirical research has also examined these issues. Johnson and colleagues interviewed 19 employees across seven United States (US) cities and described post-deployment evaluation responsibilities, including cases where buyers devised their own metrics when vendor evaluations were unavailable (Johnson et al. 2025, sec. 5.3). Kuehnert and colleagues combined stakeholder interviews and analysis of 39 GovAI FactSheets; their study distin-

guishes disclosure as a basis for dialogue from evidence adequate for technical evaluation (Kuehnert et al. 2026, secs. 3–5). Ramsay and colleagues examined procurement and early deployment in National Health Service (NHS) chest diagnostics using interviews, observations and documents, identifying the work required for integration, governance and staff engagement (Ramsay et al. 2025). These studies provide empirical precedents rather than only proposals for contract design. The present review does not replace their fieldwork or independently replicate their findings.

Cusenza provides a particularly close comparison. The available 66-page manuscript examines state contracts and subsequent amendments or renewals selected through links to litigation, civil-society complaints or allegations of harmful algorithmic use. Its selection discussion acknowledges the risk of overrepresenting deficient acquisitions. The analysis includes testing, acceptance, post-implementation monitoring and defect-management provisions. Its opening Idaho example also follows a dispute over proprietary documentation and vendor non-renewal (Cusenza 2026, 2–4, 32–33, 48–50). Its attention to operational practice makes it a close comparator. The present review adds a cross-record distinction between the origin of an assertion, the governance activity it describes and the endpoint it can support. It does not discover the general policy-to-practice problem, introduce a general-purpose assurance system, or claim that earlier contract studies lack operational examples.

Abdulbari's contractual-telemetry proposal likewise connects administrative accountability to inspectable technical evidence and correction arrangements (Abdulbari 2026). The U.S. National Institute of Standards and Technology (NIST) monitoring report distinguishes operational, functional, human-factor, security, compliance and wider-impact concerns and identifies practical barriers to monitoring deployed systems (Rao et al. 2026). A further close comparator is Hardy's 2026 Oversight-to-Control (O2C) study, which uses a purposive comparative documentary analysis of high-impact AI to distinguish procedural human involvement from effective control (Hardy 2026). Its publisher abstract reports 19 public cases and 41 oversight-control statements. This precedent is not a procurement-contract event-linkage study, but it makes a generic priority claim for documentary oversight coding untenable. This comparison relies on publisher metadata and the abstract; it is neither a full-text methods audit nor an independent reproduction. These precedents reinforce the need to avoid replacing one endpoint with another: an uptime report, a privacy-compliance correction and a reduction in service harm are not equivalent observations.

Argument-based assurance provides another methodological precedent. Guidance from the Information Commissioner's Office (ICO) explicitly links claims, contextual arguments and evidence, including for procurement and independent assessment (Office n.d.). Sabuncuoglu, Burr and Maple operationalise dynamic evidence collection through software-supported fairness assurance (Sabuncuoglu et al. 2025). Consequently, neither a source-linked register nor deterministic checks on its fields establish methodological novelty by themselves. A well-specified assurance case can represent the same qualifications used here. This article instead offers a critical documentary application and synthesis; it does not estimate an advantage over Claims–Arguments–Evidence (CAE), validate a new classifier or claim first-in-field status.

3. Review design and evidence treatment

3.1. A purposive critical review, not a completed systematic review

The primary source set consists of a US Government Accountability Office (GAO) acquisition audit, Victorian and Queensland audit reports, and an Ontario public-health implementation-team report (Office 2026,?, 2025; Chen et al. 2025). The reports were selected retrospectively for their contrasting forms of evidence: acquisition constraints and

lifecycle states; independently tested supplier changes and statutory correction; recurring performance management; and local workflow experience. The original register contains 28 event descriptions nested within nine programme families. These figures describe the retained evidence register, not nine post-deployment cases. Operational records include Victoria, Queensland and the Ontario pilot; GAO contributes historical feedback and lifecycle accounts as well as pre-award and contract-design comparators. Multiple statements from one report or programme are not independent replications. Table 1 identifies analytical roles before any substantive comparison.

The United Kingdom (UK) extension is exploratory and purposive. Seven primary records were retained to contrast two acquisition-boundary controls, one agreement with identifiable terms, a council account of collaborative supplier changes, an inspection account of staff experience, and a different-buyer procurement control. The extension retains 13 earlier entries and separates one combined Kingston entry into two attributed statements, giving 14 current interpretations. This is a correction of statement granularity, not another empirical event, contract or study. Two additional, retrospectively selected checks use a council assessment-timeline source and four official Newcastle register/company records. These records are additional sources for cross-record interpretation, not additions to the seven-document UK source set or to any denominator of government contracts. Table 1 separates these collections and their roles. Neither selection was preregistered. The UK source set was assembled after some illustrative cases were already known, and the two extra probes were chosen because they exposed linkage problems; they are not an independent holdout. There was no exhaustive search of all UK contract awards or AI disclosures. Selecting available sources does not quantify selection bias or support population-level rates. The analysis cannot estimate how often government AI suppliers are monitored, how many remedies succeed, or the prevalence of disclosure errors.

Table 1. Analytical roles of the retained documentary material.

Material and role	Evidence used in the comparison	Restriction
Operational oversight and response accounts	Victoria and Queensland audits; Kingston council account; historical FEMA geospatial feedback	Actual or attributed activity, not a common intervention, success rate or independent effect sample
Implementation and experience contrasts	Ontario pilot; North Northamptonshire staff accounts in CQC inspection	Buyer adaptation, implementation measures or staff experience; not automatically supplier-remedy outcomes
Acquisition and contractual provisions	North Northamptonshire notice and agreement; Bexley trial notice	Status and published provisions; no inference that a right was exercised or a trial completed successfully
Acquisition-scope controls	West Berkshire in-house tool; Warwickshire non-commercial outputs	Outside the purchased-AI subset; retained only to test scope interpretation
Retrospective cross-record checks	North Northamptonshire assessment period versus a later contract; Newcastle company identifier versus award contractor	Documentary mismatches, not a population sample, misconduct finding or evidence of an exercised remedy
Predeployment, design and lifecycle comparators	GAO records on PARC, SoKAT, Maven, XM-30 and USAi; separate geospatial retirement account	Do not enter a denominator of deployed supplier corrections; exact stage can remain uncertain
Selected scholarship and assurance guidance	Contract analysis, city fieldwork, disclosure research, healthcare implementation and assurance	Contextual comparison; not a comprehensively screened literature sample

These sources illustrate contrasting evidence profiles; they were not sampled to estimate procurement outcomes. All retained primary documents had a public institutional origin and a passage concerning acquisition status, oversight, a response or an operational account. The seven UK documents were selected from earlier exploration because they address these specific contrasts. The separate two-case cross-record check was chosen after the earlier findings were known; it illustrates how proposed documentary joins can fail, rather than serving as an independent holdout or error-rate estimate. This is an outcome-informed, retrospective selection rule and may overrepresent visible problems or favourable case narratives. No prospective search protocol, exhaustive source universe or saturation claim is asserted. Supplementary Tables S1–S8 document source roles, selection reasons, coding boundaries and shared-evidence dependencies. The recheck log records observed retrieval routes and failures; it does not reconstruct an undocumented historical search process.

A historical bibliographic discovery register is described separately in the Supplementary Material. It was not fully screened for this article and provides neither the inclusion denominator nor evidence of comprehensive coverage. Narrative-review reporting guidance informs transparent reporting, but does not convert purposive selection into a systematic review (Baethge et al. 2019).

3.2. Source access, extraction and independence

Each assertion retains a source identity, publisher, locator, public actor, relevant date, supported statement and prohibited inference. Where available, the record distinguishes the date an event was observed or evaluated, the publication date of the reporting document and the effective date of the relevant contract. The original event CSV is preserved unchanged. The supplementary register records the UK interpretations separately, with a parent-entry crosswalk for the Kingston split. A separate role-and-attribution table covers every retained entry. It distinguishes document author, attributed speaker, action, endpoint and source-access basis; these fields are interpretations, not independently validated labels. Missing information remains missing; a later account from another authority is not used to fill it.

Primary publication locations were revisited on 4 November 2026. A targeted source reread identified two linkage risks: a March 2026 inspection publication concerning an assessment that ended with a July 2025 site visit, before a September 2025 contract extension; and conflicting Companies House identifiers between an AI-use disclosure and a related parent-service award. These records are considered separately from the 42 retained interpretations and do not estimate the prevalence of incorrect linkages. The audits and several UK records were available as official HTML. For a subset of source PDFs, only indexed text could be checked; raw files or page images were not independently authenticated in this review. Such passages are treated as accessible published text, not as visually certified signatures or legal opinions. The Ontario report is an implementation-team account, not an independent audit, and its workflow findings have not been independently replicated by the other selected programmes. Individual source-access qualifications and the precise extraction history are documented in the supplement.

Source authority is coded at the statement level. An independent auditor's own finding differs from an agency statement reported by that auditor. Likewise, independent publication of a staff interview does not turn the interview into an independently measured effect. Evidence supporting supplier identity is also different from evidence supporting the performance of its product. Without these distinctions, accumulating documents can appear to provide independent corroboration of what remains a single self-reported observation.

Extraction, source interpretation and drafting were AI-assisted. No independent human coding, blinded adjudication or inter-rater agreement is claimed. The accompanying code checks identifiers, source links and preservation of the historical record; it cannot decide source truth, contractual enforceability or causal validity. The targeted rereading preserved each speaker, action, unit and relevant limit on inference. In Kingston, for example, it separated staff-reported writing-time savings from a project-reported transcription-accuracy figure. Earlier annotations remain accessible and were not retrospectively relabelled as independently validated.

3.3. Non-ordinal dimensions rather than a ladder of success

The analysis uses six separate questions, shown in Table 2. The questions do not form a maturity score. A statutory response can be documented despite inaccessible contractual terms; a measured workflow difference can exist without a supplier repair; and an executed agreement can exist before deployment. These are different evidence profiles, not positions on a universal ladder.

Table 2. Questions used to distinguish evidence without imposing an ordinal score.

Dimension	Question	Example of an invalid inference
Acquisition	What was acquired, if anything, and at what stage?	An algorithmic disclosure implies a purchased AI tool.
Identity and scope	Do buyer, supplier **legal ID**, product, service and period match?	A matching brand or a common platform proves the same legal contractor and AI module.
Authority	What power or process underlies the action?	Statutory scrutiny proves exercise of a contractual cure clause.
Action and object	Who changed what, and when?	Staff template changes prove a supplier model release.
Measurement	What endpoint, unit, **observation period**, publication date and assessor are recorded?	A report published after a contract describes performance during that contract.
Attribution	What links the action to the claimed result?	A later favourable metric is the effect of a contractual remedy.

The contractual entitlement–exercise–supplier action–measurement–attribution sequence applies to the specific compound claim that exercising a contract right caused a measured benefit. Such a claim requires applicable authority, a linked exercise, a supplier response and a credible outcome evaluation. This is not a requirement for every claim about accountability. The requirements for a narrower assertion are limited to the proposition actually being made.

4. Findings

4.1. A government algorithm register is not an AI-contract register

The West Berkshire disclosure describes a rules-based assessment tool created internally on a commercial platform and explicitly reports no procurement specific to that tool (Council 2025). Warwickshire reports non-commercial provision of model outputs from a staff-originated personal project associated with an external company (Council 2025). Both are informative scope controls, not observations of failure to enforce a purchased-AI contract.

The acquisition distinction determines which documents belong in an analysis of supplier accountability. Excluding an internally built tool from a purchased-AI subset does not exclude it from algorithmic-governance research. Similarly, a commercial platform can underpin an in-house application without proving that the application's decision logic was

purchased from that platform provider. Supplier, platform, model and service roles must therefore be stated rather than inferred from a company name.

The seven-source extension is defined by its explicit source list, not by an assumed census of local-authority disclosures or procurement contracts. Directory tags and retrieval versions do not establish a stable procurement sampling frame; the resulting coverage uncertainty is recorded in the supplement.

4.2. Two cross-record linkage checks: observation time and legal entity

A document's publication date need not coincide with the period of experience it describes. The Care Quality Commission (CQC) published its North Northamptonshire report in March 2026, and that report contains staff testimony about the Magic Notes pilot (Commission 2026). The council dates the assessment from February 2025 through an on-site visit in July 2025 (Council 2026). The specific twelve-month extension examined here was signed on 3 September 2025 (Council 2025). The individual staff statements are not separately dated, so the report's March 2026 publication cannot establish that they describe performance under the September extension or any exercise of its terms. An earlier pilot agreement or later use may have existed, but its contractual implications cannot be deduced from the publication date.

A distinct legal-identity issue appears in Newcastle. An official Algorithmic Transparency Recording Standard (ATRS) disclosure for Amazon Q and Contact Lens names "AWS" with Companies House number **09162820**, whereas the council's official award for its Amazon Connect contact-centre service identifies the contractor as **AWS EMEA SARL UK**, company **FC034225** (Council 2026, 2025). Companies House associates 09162820 with AWS (UK) LIMITED, an engineering-activities company, and FC034225 with AMAZON WEB SERVICES EMEA SARL (House 2026,?). The buyer and parent platform are consistent at the description level, but the published legal identifiers do not match. The award does not separately establish the exact purchase terms for each named AI module. This is an apparent record-level entity-identifier inconsistency, **not** a legal finding of misrepresentation, deliberate error, breach or contract failure. Table 3 presents these two distinct risks of erroneous linkage alongside the previously documented cross-buyer control.

Table 3. Cross-record linkage checks that block tempting but unsupported conclusions.

Proposed join or interpretation	Anchored observation	Permitted conclusion and hold
North Northamptonshire CQC publication (March 2026) → September 2025 extension performance	Council dates assessment February–July 2025; twelve-month extension signed 3 September 2025; exact dates of individual staff accounts unavailable.	Staff testimony is retained, but its publication date does not prove a post-extension effect or exercised remedy.
Newcastle ATRS "AWS" → award contractor by brand name	ATRS legal ID 09162820; related parent-service award ID FC034225; these resolve to distinct registered entities.	Same buyer/platform context, **legal-entity identity unresolved across the records**; exact AI-component contracting and wrongdoing not established.
Bexley Magic Notes → North Northamptonshire remedy terms by shared supplier	Two different authorities name Beam Up Ltd; Bexley records a 2026 trial, NNC a 2025 extension.	Supplier/product overlap does not transfer contractual rights or outcomes across buyers.

These retrospective checks identify unsupported joins; they do not estimate how often such problems occur. In the first case, chronology blocks an apparent post-contract evaluation; in the second, identity reconciliation blocks a false exact-company match; in the third, separate public buyers block terms transfer. None requires concluding that an unknown underlying event never occurred. The two new probes have a separate source register and are not counted as independent additional effectiveness studies.

4.3. An executed agreement establishes provisions, not their exercise

North Northamptonshire's official notice reports a £48,000 Magic Notes agreement signed on 3 September 2025, with a term from 8 September 2025 to 7 September 2026, and links the executed instrument (Council 2025). Its supplier identifier matches the company identified in the instrument. This supports a narrow buyer–supplier–product linkage.

The published agreement text contains three different provisions (Council and Ltd 2025, 3, 7, 16–17): a conditional reasonable-commercial-endeavours undertaking to correct nonconformance with reasonable skill and care (not an error-free guarantee); a data-processing annex allowing specified compliance audits subject to purpose and notice conditions; and a general-terms sentence restricting access to the product for monitoring its availability, performance or functionality, or for benchmarking. That restriction alone does not show that statutory powers or every form of internal review were suspended. Conversely, the processing-compliance inspection provision does not automatically confer general model-benchmarking permission. This is an *audit-object and access-scope asymmetry* visible in the published text, not an adjudicated legal conflict or a finding that one clause defeats the other. Resolving their interaction, applicable duties and any exceptions requires legal and contractual information beyond this review.

The order also specifies 11,900 recording hours, 100 reports and usage reporting (Council and Ltd 2025, 1–2). These are subscription and reporting commitments, not demonstrated consumption or staff time saved. No dated exercise of the correction undertaking was retrieved. Table 4 therefore separates this provision-level evidence from observed responses in other sources.

Table 4. Distinct evidence profiles; rows are not successively stronger levels.

Setting and source	Supported documentary observation	Claim that remains unsupported
North Northamptonshire agreement	Published contractual terms, with distinct correction, processing-inspection and access provisions	That a particular term compelled an observed supplier correction
Victoria audit	Tested, buyer-approved software rollout; separate audit-linked archiving correction	That contractual enforcement caused improved model accuracy or road safety
Queensland audit	Recurring audits, performance reviews and feedback to the vendor	A specified release and an isolated post-release effect
Kingston council account	Reported collaborative supplier adjustments and continuing staff feedback	A clause-specific compelled repair or independently estimated remedy effect
Ontario implementation account	Reported buyer adaptation and divergent local workflow measures	A supplier-remedy causal effect
North Northamptonshire CQC account	Staff experience recorded within an independent inspection publication	Independent quantitative validation of efficiency or throughput

4.4. Documented oversight and supplier responses are not necessarily contractual sanctions

The Victoria audit supports two different supplier actions. A buyer-approved software version was rolled out after independent testing in December 2025; its performance effect was not yet established. Separately, the auditor linked correction of image archiving to the audit and reported that the resulting retention arrangements complied with the relevant regulations (Office 2026). The seven-year archive retention period exceeded the applicable regulatory limits.

The latter is positive documentary evidence, within the auditor's published account, of a supplier response associated with statutory scrutiny. It does not establish a model-

accuracy improvement or reveal the exercise of a particular contractual clause. Recognising that response is compatible with withholding those stronger claims. This distinguishes a documented governance response from a causal evaluation of contract effectiveness.

Queensland provides a different profile: contractual performance indicators, weekly audits, quarterly reviews and an established feedback route to the external vendor (Office 2025). The auditor positively assesses these processes as supporting refinement, accuracy and fairness. That qualitative assessment should be retained, not overwritten by a claim that oversight was ineffective. The report's workflow-filtering measure is nevertheless not a dated, isolated patch-effect estimate. A favourable operational statistic and recurring governance activity can coexist without identifying which supplier change, if any, caused the statistic. Routine oversight should therefore be described as documented activity rather than discounted as meaningless because no sanction is visible.

4.5. Collaborative changes, human adaptation and favourable reports are separate observations

Kingston's case narrative describes continued feedback and supplier adjustments addressing technical difficulties (Association and Council 2025). Its favourable measures have different origins: the 50–60% reduction in note-writing time is attributed to social workers, whereas transcription accuracy above 96% is presented as a project result. The account does not provide the denominator, scoring rule or release-specific comparison needed to independently evaluate the latter. Neither is an independently estimated effect of a contractual cure. These statements are therefore retained as distinct attributed claims, not a single staff-reported performance measure. Staff also reportedly amend outputs and refine templates. This makes attribution to the supplier alone especially uncertain: the service combines software, professional review, configuration and changing working practices.

The separate Bexley notice describes a £30,000 three-month Magic Notes trial signed on 1 February 2026 (Bexley 2026). It is a useful different-buyer control, not an independent replication of Kingston's reported outcomes. Neither Bexley's nor Kingston's evidence can acquire North Northamptonshire's contractual terms merely because the product and supplier names match.

The Ontario account illustrates the same attribution difficulty in another setting. Its report describes opposite-direction documentation-time differences at two public-health units, together with local template work after initial vendor support (Chen et al. 2025). The report-level total of 44 interactions is retained descriptively; a conflicting narrative count of 42 is noted in the supplementary access record. No denominator is reconstructed, no sites are pooled and no new effect estimate is calculated. The inconsistency affects quantitative precision; it does not remove the source's qualitative account of local template work. The lack of fresh original-page verification further limits the evidentiary weight of this example. Its contribution is the warning that an implementation outcome need not be a supplier-remediation outcome.

4.6. Independence belongs to an observation, not merely to its publisher

CQC's North Northamptonshire assessment records staff reports of clearer, more consistent and efficient records using Magic Notes (Commission 2026). The council dates the assessment to February–July 2025, before signature of the September 2025 extension; the statements themselves have no separate interview dates. Publication in March 2026 therefore does not establish testimony during that later contract (Council 2025, 2026). CQC is an independent inspection publisher, but the relevant statements remain attributed staff experience. They cannot be used as independent numerical validation of council efficiency

claims. The same distinction applies when an audit repeats an agency's account: the audit may independently verify some findings while attributing others to officials.

GAO's procurement audit further shows why programme and supplier states should not be collapsed. Its accounts of broader vendor use and a specific AI programme's retirement concern different objects and are not sufficient to establish a common contractual chronology. Pre-award licensing decisions, follow-on requirements and retirement before production also belong outside a denominator of post-deployment supplier cures (Office 2026). These contextual records help define the limits of the main question rather than expanding an apparent success-or-failure sample.

Across the selected material, a causal estimate linking an exercised contractual remedy to a public-service outcome remains unsupported. This is a bounded finding about accessible evidence, not a finding that no remedies occurred, that contracts failed generally, or that the observed governance activities had no value.

4.7. Dependence and alternative interpretations

The supplier-rollout and retention-correction accounts both come from the Victoria audit and programme. They are two differently scoped actions, not independent replications of a general accountability effect. Likewise, the UK commercial examples concentrate on the same product family and supplier; multiple public buyers do not establish supplier diversity. The provision-level reading depends on the North Northamptonshire agreement, and the collaborative-response account depends on the Kingston narrative. The source-dependency map in the supplement records these relationships without converting coverage into an accuracy or robustness percentage.

At least three explanations remain compatible with an absent public remedy notice: no relevant defect arose; an issue was resolved through informal or private arrangements; or a formal remedy was used but not published. The selected records cannot distinguish these alternatives. Conversely, a favourable account may reflect the AI tool, human review, training, configuration changes or several of them together. This review identifies the documented observation and its source, not which unobserved explanation is correct.

5. Discussion

5.1. Recognising activity without overstating effects

A review can report documented practices without estimating their effects, but should not rename those practices as achieved accountability. The two distinct cross-record checks illustrate this narrower message: matching the buyer and brand is insufficient if legal identifiers conflict, while a report published after an award does not itself date its underlying observations to the contract period. The comparison therefore requires keeping these distinct claims separate. The present synthesis distinguishes provision-level evidence, recurring oversight, source-linked supplier response, reported operational experience and causal evaluation. These distinctions are established in assurance and evaluation research. Applying them here shows why apparently conflicting observations can coexist. Hardy's O2C analysis independently underscores the wider point that nominal oversight and effective control differ (Hardy 2026). That broader insight belongs to prior work. The examples in this review concern a contract's specified audit object, the identity of its legal parties and the dates of underlying observations relative to publication. City fieldwork on buyer-led evaluation (Johnson et al. 2025), disclosure research on the limits of vendor self-reports (Kuehnert et al. 2026) and implementation research on organisational work (Ramsay et al. 2025) provide complementary context. Their observations support attention to actors and information provenance; they do not validate the present case interpretations or turn them into independent replications.

An agency may have a valid contract, conduct regular reviews, receive a supplier adjustment and still lack a credible estimate of the adjustment's effect. A different agency may obtain a compliance correction through statutory scrutiny without a public cure notice. A third may improve its workflow through staff adaptation rather than a supplier release. None of these profiles can be represented faithfully by a single binary question about whether a complete contractual chain was established.

This interpretation does not weaken the evidence requirement for causal claims. It assigns the requirement to the correct claim. To estimate the effect of a remedy, researchers would still need the intervention's scope and timing, a measurement plan and an appropriate counterfactual or other defensible identification design. Placing a contract, a later press account and a favourable percentage in chronological order does not supply those data.

5.2. From a clause inventory to a usable public record

The practical implication is to specify the records needed when oversight occurs, rather than merely recommending additional clauses. Table 5 sets out a claim-to-record checklist proposed from the synthesis and the established monitoring and assurance literature (Rao et al. 2026; Office n.d.). It is a documentation recommendation, not a validated standard or proof that publication alone improves outcomes. It can be incorporated into an existing assurance case rather than competing with one.

Table 5. Proposed claim-to-record checklist, not a validated standard.

Record to preserve	Suggested content	Supported purpose
Acquisition and scope record	Buyer, supplier **registered legal ID**, product/service and any AI-module scope, contract reference, effective period and variations	Prevent cross-buyer, cross-entity, cross-component and cross-period linkage
Oversight episode	Issue, date, authority or collaborative process, responsible actor and evidence requested	Distinguish a provision from its actual use
Supplier response	Dated change, affected component/version, scope, acceptance or unresolved issue	Identify what changed without assuming benefit
Measurement record	Endpoint, numerator/denominator, **observation period separately from publication date**, assessor, human workflow and concurrent interventions	Interpret performance without backdating or postdating an outcome to a contract
Evaluation statement	Claim, comparison design, assumptions and remaining uncertainty	Separate operational reporting from causal attribution

Public release need not include personal records, raw service-user audio, proprietary source code or unrestricted access to a deployed system. A public summary can identify the event, scope and verification method while protected evidence is retained for authorised assessment. This arrangement is consistent with established assurance approaches that distinguish public claims from controlled access to supporting evidence (Office n.d.). The exact disclosure and testing permissions require case-specific legal and information-governance review.

5.3. What this review adds—and does not add

The empirical contribution lies in comparing evidence objects often grouped under a single accountability label. The chronology and legal-entity checks are concrete, source-backed illustrations of unsupported joins—not novel matching algorithms or a statistically

representative error survey. Scope controls show why algorithmic transparency and procurement inventories are not equivalent. The agreement example offers a source-checkable contrast between restricted performance-monitoring access and a specified data-processing-compliance inspection right, without adjudicating enforceability or assuming that the buyer lacked every monitoring power. The response examples separate statutory correction, collaborative improvement and routine performance management. The inspection example separates institutional independence from the measurement status of a reported outcome.

The contribution is a comparative account of specific acquisition, timing, identity and attribution problems, including different evidence origins within the same publication. It is not a new theory of accountability or evidence that the present coding outperforms a competent reviewer or a CAE implementation. No sensitivity, specificity, accuracy or comparative-effectiveness measure is estimated. The accompanying structured records make this review easier to inspect; they do not establish general superiority. The closest scholarship already addresses contract design, operational disputes and evidence-producing arrangements (Cusenza 2026; Abdulbari 2026). The defensible contribution is an incremental, transparent synthesis of selected records and their implications for documenting oversight. No claim is made that the observed evidence profiles form an exhaustive taxonomy.

5.4. Limitations and a focused research agenda

Selection is the principal limitation. Sources were chosen for informative contrasts and availability, and some were discovered after the initial framing. Public audits may overrepresent problematic or consequential cases, while council case studies may emphasise favourable experiences. Neither source class estimates what happens across all public AI acquisitions. The UK extension is concentrated on Magic Notes for commercial examples, and two Victorian supplier actions share one audit and programme. Sector, supplier and source concentration therefore constrain transferability even when the individual passages are correctly interpreted. UK directory tags and retrieval versions add coverage uncertainty. No geographical or supplier ranking is justified.

Access and interpretation also limit the conclusions. Fresh page-image checks did not succeed for the Ontario report, North Northamptonshire agreement, Bexley notice or full CQC report, despite accessible indexed official text. Private correspondence, issue logs and amendments may change the interpretation of an episode. The absence of a public exercise record cannot distinguish a privately resolved issue from no issue, no exercise or missing disclosure. Source assertions and code are AI-assisted and remain susceptible to semantic error despite structural checks.

Future work should address distinct questions separately. To study contract-enforcement effects, obtain dated notices, supplier release records, acceptance evidence and an evaluation dataset from the same buyer and programme. To test documentary coding performance, use independently adjudicated source passages and a matched assurance comparator. To estimate prevalence, define a stable sampling frame that includes acquisition status and document availability. These questions require different study designs. Not completing one does not imply that another has been completed, nor does it invalidate the narrower documentary findings reported here.

6. Conclusion

Public evidence can document oversight and supplier responses without establishing achieved accountability or a causal contract-enforcement effect. The selected records describe recurring oversight, supplier deployment, a statutory-audit-linked correction and reported collaborative changes, while exposing acquisition, legal identity, observation chronology, authority and measurement boundaries. These activities should be assessed

on their own evidentiary terms, neither merged into a single success chain nor dismissed because stronger causal evidence is unavailable. A useful next step for public buyers is to preserve linked acquisition, oversight, response and measurement records with explicit scope and source provenance. For researchers, the central obligation is to match the strength of the claim to the evidence that actually supports it.

7. Data and code availability and AI-use disclosure

The Supplementary Material documents source selection, attributed interpretations, a parent-entry crosswalk, two retrospective cross-record checks, analytical roles and access qualifications. A public subset of the retained event registers and deterministic checks is deposited at https://github.com/shaikhamalkawi-ux/governing-procured-ai-review. The internal versioned archive contains additional source and provenance material that is not publicly redistributed, including licensed bibliographic exports and copyrighted scholarly PDFs. The public structural replay verifies consistency of the 28 event descriptions and 140 proposition checks against frozen, AI-assisted interpretations; it does not establish legal meaning, independent source truth or causal attribution. The repository also includes a separate, incomplete systematic scoping-review checkpoint whose screening labels are not final eligibility decisions or the documentary sample. No permanent archive DOI has been assigned. No personal participant-level data were analysed. Generative AI assisted literature discovery, source interpretation, code construction and drafting; no independently adjudicated coding or human inter-rater statistic is claimed.

Supplementary Materials: Supplementary Material, including Tables S1–S8 and document-level evidence registers, accompanies the manuscript as a separate file.

Data Availability Statement: Public, non-licensed research materials and executable structural checking scripts are available at https://github.com/shaikhamalkawi-ux/governing-procured-ai-review. The repository includes the frozen 28-event documentary source-to-claim register and its 140 deterministic proposition checks. Passing this replay verifies rules and retained inputs, not independent truth of the original source interpretation. It also holds a separate, explicitly provisional scoping-screening checkpoint (220 of 2556 bibliographic identities, without final eligibility or PRISMA-ScR claims). Government source documents are linked at their official publication sites; licensed Scopus and Web of Science exports and third-party copyrighted full texts are not redistributed. No permanent archive DOI has been assigned.

Acknowledgments: Generative AI assisted literature discovery, source interpretation, code construction, and manuscript drafting. No independent human source adjudication was performed.

References

Coglianese, Cary and Lampmann, Erik. 2021. Contracting for Algorithmic Accountability. *Administrative Law Review Accord*, vol. 6, pp. 175–199. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3853208.

UK Government. 2020. Guidelines for AI Procurement. https://www.gov.uk/government/publications/guidelines-for-ai-procurement/guidelines-for-ai-procurement.

Sánchez-Graells, Albert. 2024. Public Procurement of Artificial Intelligence: Recent Developments and Remaining Challenges in EU Law. *Legal Tech Journal*, vol. 2, pp. 122–131. https://doi.org/10.2139/ssrn.4706400.

Cusenza, Giulia G.. 2026. From Policy to Practice: Closing the Gap in AI Public Procurement. *Drexel Law Review (forthcoming); SSRN working manuscript*, vol. 19. 66-page manuscript consulted; publication metadata provisional. https://doi.org/10.2139/ssrn.6812838.

Abdulbari, M.. 2026. The Contractual Vacuum: AI Procurement and the Erosion of Good Administration. *Frontiers in Artificial Intelligence*, vol. 9, pp. 1916534. https://doi.org/10.3389/frai.2026.1916534.

Rao, A. and Keller, A. and Kalra, N. and Steed, R. and Kwegyir-Aggrey, K. and Klyman, K. and Staheli, D. and Bergman, A.. 2026. Challenges to the Monitoring of Deployed AI Systems: Center for AI Standards and Innovation. *NIST AI 800-4*. https://doi.org/10.6028/NIST.AI.800-4.

Information Commissioner's Office. n.d.. Annexe 5: Argument-Based Assurance Cases. Undated online guidance, under review at access; publication year unverified. https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/artificial-intelligence/explaining-decisions-made-with-artificial-intelligence/annexe-5-argument-based-assurance-cases/.

Sabuncuoglu, Alpay and Burr, Christopher and Maple, Carsten. 2025. Justified Evidence Collection for Argument-based AI Fairness Assurance. *FAccT 2025*, pp. 18–28. https://doi.org/10.1145/3715275.3732003.

U.S. Government Accountability Office. 2026. Artificial Intelligence Acquisitions: Agencies Should Collect and Apply Lessons Learned to Improve Future Procurements. *GAO-26-107859*. https://files.gao.gov/reports/GAO-26-107859/index.html.

Victorian Auditor-General's Office. 2026. Reducing the Harm Caused by Distracted Drivers. https://www.audit.vic.gov.au/report/reducing-harm-caused-distracted-drivers.

Queensland Audit Office. 2025. Managing the Ethical Risks of Artificial Intelligence. *Report 2: 2025–26*. https://www.qao.qld.gov.au/reports-resources/reports-parliament/managing-ethical-risks-artificial-intelligence.

Chen, H. and Hirschfeld, C. and Shah, R. and Shakeel, N. and Rebellato, S. and Riewe, N. and others. 2025. Enhancing Public Health Disease Investigations through AI Scribe Technology: A Pilot Study in Two Ontario Public Health Units. *Public Health Ontario-hosted project report*. Source text available; fresh original page-image verification not completed. https://www.publichealthontario.ca/-/media/Documents/L/26/ldcp/ldcp-disease-investigations-ai-scribe-technology.pdf?hash=C5741362E1103668EFEC6345082465E4&rev=11724f68f84a449bad7e7a015b40ef53&sc_lang=en.

Baethge, C. and Goldbeck-Wood, S. and Mertens, S.. 2019. SANRA—A Scale for the Quality Assessment of Narrative Review Articles. *Research Integrity and Peer Review*, vol. 4, pp. 5. https://doi.org/10.1186/s41073-019-0064-8.

West Berkshire Council. 2025. Apply for a Larger Rubbish Bin: Algorithmic Transparency Record. *GOV.UK*. https://www.gov.uk/algorithmic-transparency-records/west-berkshire-council-apply-for-a-larger-rubbish-bin.

Warwickshire County Council. 2025. Domestic EPC Estimates: Algorithmic Transparency Record. *GOV.UK*. https://www.gov.uk/algorithmic-transparency-records/warwickshire-county-council-domestic-epc-estimates.

North Northamptonshire Council. 2025. Magic Notes: 12 Month Contract Extension of Pilot, UK7 Notice 2025/S 000-053743. *Find a Tender*. https://www.find-tender.service.gov.uk/Notice/053743-2025.

North Northamptonshire Council and Beam Up Ltd. 2025. Magic Notes Subscription: Order Form and General Terms of Service, Official Attachment A-4800. *Find a Tender*. Source text available; fresh original page-image verification not completed. https://www.find-tender.service.gov.uk/Notice/Attachment/A-4800.

Local Government Association and Kingston Council. 2025. Kingston Council: Using AI in Adult Social Care Administration. *Local Government Association*. https://www2.local.gov.uk/case-studies/kingston-council-using-ai-adult-social-care-administration.

London Borough of Bexley. 2026. Magic Notes Solutions, UK7 Notice 2026/S 000-018512. *Find a Tender*. https://www.find-tender.service.gov.uk/Notice/018512-2026.

Care Quality Commission. 2026. North Northamptonshire Local Authority Assessment: Learning, Improvement and Innovation. *CQC*. https://www.cqc.org.uk/care-services/local-authority-assessment-reports/northnorthamptonshire-0326/theme4/learning.

Johnson, N. and Silva, E. and Leon, H. and Eslami, M. and Schwanke, B. and Dotan, R. and Heidari, H.. 2025. Legacy Procurement Practices Shape How U.S. Cities Govern AI: Understanding Government Employees' Practices, Challenges, and Needs. *FAccT 2025*, pp. 772–789. https://doi.org/10.1145/3715275.3732049.

Kuehnert, Blaine and Johnson, Nari and Dotan, Ravit and Heidari, Hoda. 2026. Disclosure or Marketing? Analyzing the Efficacy of Vendor Self-reports for Vetting Public-sector AI. pp. 687–

713. Preprint arXiv:2604.01332v1 consulted for method; proceedings identity verified separately. https://doi.org/10.1145/3805689.3812220.

Ramsay, A. I. G. and Crellin, N. and Lawrence, R. and others. 2025. Procurement and Early Deployment of Artificial Intelligence Tools for Chest Diagnostics in NHS Services in England: A Rapid Mixed-method Evaluation. *eClinicalMedicine*, vol. 89, pp. 103481. https://doi.org/10.1016/j.eclinm.2025.103481.

North Northamptonshire Council. 2026. Health and Adult Social Care Scrutiny Committee: Adult Social Care CQC Assessment Outcome and Improvement Plan, Minutes, 2 July 2026. *Council minutes*. https://northnorthants.moderngov.co.uk/mgAi.aspx?ID=16569.

Newcastle City Council. 2026. AWS Contact Centre Services (Amazon Q and Contact Lens): Algorithmic Transparency Record v4.0. *GOV.UK*. https://www.gov.uk/algorithmic-transparency-records/newcastle-city-council-aws-contact-centre-services-amazon-q-and-contact-lens.

Newcastle City Council. 2025. Customer Contact Centre through Digital and Omni-channel, F03 Award 2025/S 000-017533. *Find a Tender*. https://www.find-tender.service.gov.uk/Notice/017533-2025.

UK Companies House. 2026. AWS (UK) LIMITED: Company Number 09162820. *Companies House*. https://find-and-update.company-information.service.gov.uk/company/09162820.

UK Companies House. 2026. AMAZON WEB SERVICES EMEA SARL: Company Number FC034225. *Companies House*. https://find-and-update.company-information.service.gov.uk/company/FC034225.

Hardy, Karim. 2026. From Human Oversight to Effective Control: A Comparative Documentary Framework for High-Impact AI Systems. Publisher abstract and metadata checked; full text not independently evaluated. https://www.accscience.com/journal/DP/articles/online_first/9411.